Using Sentiment Analysis to Detect Hate Speech

By: Varun Chilukuri, Rahul Nair, Anirudh Porur, Shreya Shete

Problem

- Millions of people use social media platforms like Youtube and Twitter every day
- Such platforms oftentimes have hateful and offensive language, which can influence impressionable youth.

Problem (cont.)

Project Goals:

- To identify hateful/offensive comments so they can be removed on these platforms
- Develop a user-friendly tool that utilizes our model's results to classify comments

What makes this challenging?

- Hate speech can be subtle and heavily context-dependent
- Language is constantly evolving and new phrases are being coined all the time
- We need to balance sensitivity and specificity—overly aggressive models may result in high false positives, impacting free speech and causing user dissatisfaction.

Dataset

Hate Speech and Offensive Language Dataset from Kaggle:

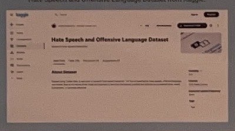

Word Cloud

"Neither Hate Speech nor Offensive Language"

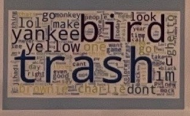

Class Distribution After Upsampling and Downsampling

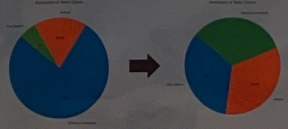

Approach / Methods

Libraries: Tensorflow, Keras, Numpy, Pandas, Nltk, Sklearn, matplotlib

Data preprocessing (for Tweet column)

- balancing dataset by upsampling hate speech & downsampling offensive speech
- removing punctuation & stop words
- splitting text into words
- tokenizing, converting to integers
- padding sequences to have same length

Bi-Directional LSTM Model

- Input: 1D Sequence of Padded Tokens
- Output: Classifies as Hate Speech and/or Offensive Language, or Neither

- Training & Testing Model
- Plotting Accuracy and Loss for Training & Validation Data over Epochs
- Calculating Accuracy Precision, Recall, and F1 Score for Test Data

Model Architecture

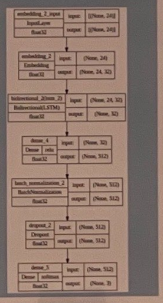

Results

```
Accuracy on test data: 0.9067562222488774
F1 Score: 0.9087254838943481
Precision: 0.9028119905471882
Recall: 0.8994792103767395
```

We achieved a validation accuracy of nearly 91%

Precision, Recall, F1 score of nearly 90% → algorithm returns more relevant results

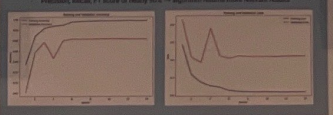

Applications

Text Classifier

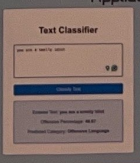

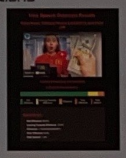

Discussion

Success Criteria

- The model accurately identifies hate speech, offensive language, and neutral speech
- We achieved a high accuracy and built two user-friendly tools that allow us to analyze individual comments or an entire YouTube video section

Roadblocks

- Initial dataset improperly labeled hate speech values as "neutral"
- New dataset had more nuance—it differentiated between "offensive language" and "hate speech"

Our Solution

Limitations

- Profanity and true hate speech is classified well
- Friendly and supportive text is classified well
- Strong real-world applications

Future Work

- Gather more data or improve on size of data with more segmentations
- Improve accuracy of model in general
- Classify mildly offensive text more accurately
- Provide a more detailed text classification rather than just "offensive language", "hate speech", or "neither"